

Dyalog APL in Non-Windows Environments

Miscellaneous features etc which you might come across when using Dyalog APL on
non-Windows platforms

Andy Shiers, COO Dyalog Ltd

Supported Platforms

- Windows: 32- and 64-bit, Classic and Unicode
 - Linux: 64-bit, Classic and Unicode
 - macOS: 64-bit Unicode
 - Raspberry Pi: 32-bit Unicode on 32-bit Raspbian
-
- AIX: 32- and 64-bit, Classic and Unicode
-
- 64-bit Unicode for Windows, Linux and macOS, and 32-bit Unicode for Raspberry Pi can be downloaded from www.dyalog.com

Major differences

- 32-/64-bit: size of array/workspace, interaction with O/S
- Classic/Unicode: Number of supported characters, font limitations etc
- Windows vs the rest: full IDE vs tty interface
 - RIDE can be used to drive recent Dyalogs on any platform
- AIX vs the rest: Big Endian vs Little Endian
- Not a huge difference between AIX/Linux&Pi/macOS
 - But don't assume that they are always the same !

Major differences – workspace

- Workspaces and component files should be compatible across all platforms
 - Within the same version
 - Are upwardly compatible
 - Classic generates TRANSLATION ERRORS if a workspace contains characters which are not in □AVU
 - Watch out for such characters in □DM and □DMX !
 - GUI features will be deleted when loading a Windows workspace into a non-Windows version
 - Will get warning messages .. ignore them!
- WSEXT/CFEXT control what extensions workspace and component file have (if any)

Major differences – workspace and array size

- Workspace size (MAXWS)
 - Maximum size to allow; must be contiguous in the process space
 - 32-bit .. Maximum is between 1.2 to <2.0GB
 - 64-bit .. "unlimited" .. biggest we know of is 2.1TB
 - Don't make too big or O/S will start to swap which really slows things down
 - Consider size of workspace × number of concurrent instances
- An array must fit in workspace:
 - 32-bit can have up to 2^{32} elements
 - 64-bit "unlimited"

Major differences – Classic/Unicode

- Classic has at most 256 characters
 - Limited further by font considerations
 - Much more flexibility with keyboard setup (APLK/APLKEYS etc)
- Unicode has considerably more
 - 1, 2 or 4 byte values
- Even in Classic better to use □UCS, not □AV/□TC

Big Endian/Little Endian

- AIX runs on POWER CPU in big-endian mode (the original mode)
- All other supported platforms are little-endian
 - Be aware that you may have to consider byte ordering
 - Intel/AMD/ARM: Little endian – 1 2 3 4
 - POWER (AIX): Big endian – 4 3 2 1
 - eg 11 □ DR 123456
- Dyalog APL does most of the work for you
 - Always)SAVEs in current endianness (!)
 - Components written in local endianness

Major differences – IDE vs tty

- Windows IDE
 - GUI based, floating resizable windows
- tty interface
 - Assumes fixed sized, character based output
 - Responds to some resize requests – might need screen refresh
- RIDE
 - Client for all but AIX
 - Can talk to any Dyalog after 16.0
 - Zero-footprint RIDE (runs in browser)
 - included with all non-Windows versions of Dyalog

Auxillary Processors

- Example: 'xutils'□sh"
 - Separate process which contains user-written c-code
 - Appears in workspace as locked functions
 - Erasing all the functions terminates the AP
 - Benefit: if an AP crashes, it doesn't bring the APL down too
 - But: data has to be passed to and from the AP .. Slow
- xutils
- nfiles
- unixfiles
- Customer written APs

□SM

- A screen manager
 - Available on Windows, but only to support existing applications
 - Available on non-Windows, tty interface only, only for existing applications
- Assign/change □SM to update the screen
 - A complex nested array which can contain sub-screen manager arrays
- Use □SR to pass control to user (and make □SM screen visible)
- Use HK keystroke to toggle □SM screen when developing

UNIX cheat sheet

- Big difference with Windows:
 - Windows: erase *.*
 - The erase command sees "*" so can ask if you really meant to do it.
 - UNIX: rm *
 - The command shell expands "*" to all that matches it (ie everything). So rm has no idea that it's everything, so won't check
 - But: rm -i
- Similar for all commands. Few have the interactive option.

UNIX cheat sheet - directories

- cd
 - cd go home
 - cd - go to previous directory
 - cd adir cd to adir
- mkdir adir
- rmdir adir delete iff it's empty
- rm -r adir recursively delete (be careful!)
- rm -rf adir forcibly delete recursively (be exceedingly careful!)

UNIX cheat sheet –dealing with files

- `ls`
 - `ls -lrt` long listing in reverse time order – current directory
 - `ls f1 f2`
 - `ls *`
 - `ls adir`
- `rm`
 - `rm f1` delete f1
 - `rm *` delete all files in current directory

UNIX cheat sheet – dealing with file content

- `cat file`
 - `cat *` concatenate all files
 - `cat andy* > out` concatenate all files whose name starts with andy into a file called out
- `wc -l file` report number of lines
- `grep andys /etc/passwd` report all lines which contain "andys" in the file /etc/passwd
- `vi myfile` Use the standard UNIX editor to edit/look at myfile

UNIX cheat sheet – the file command

- file myfile
 - Attempt to describe the contents of the file
- file -m /opt/mdyalog/17.1/64/classic/p9/magic myfile
 - As above, but try to identify details of Dyalog files (the -m only needed on AIX)

```
andys@p9-72qa:/dev/tmp/andys$ file -m /opt/mdyalog/17.1/64/classic/p9/magic *
182.dws: Dyalog APL workspace type 18 subtype 33 64-bit unicode little-endian
big.dws: Dyalog APL workspace type 17 subtype 16 64-bit classic big-endian
isfile: Dyalog APL workspace type 12 subtype 5 64-bit classic big-endian
isfile121: Dyalog APL workspace type 12 subtype 5 64-bit classic big-endian
runall: c program text File contains 8 bit characters.
types.dws: Dyalog APL workspace type 18 subtype 4 64-bit unicode big-endian
xx.dcf: Dyalog APL component file 64-bit level 1 journaled checksummed
```

UNIX cheat sheet – filesystems

- Windows "sees" separate drives, each with its own filesystem
- UNIX "sees" a single directory structure, made up of one or more filesystems
 - `df` # report space available in each filesystem
 - `df -m` # report in MB
 - `du myfile` # report space used by myfile
 - `du -ms *` # report space used by each object in current directory
- `du` and `df` both based on number of blocks used. `df` is more precise

UNIX cheat sheet – pipes, redirection, exit codes

- `mycmd <infile >outfile 2>errfile`
 - 0: stdin (usually the keyboard)
 - 1: stdout – where output goes (usually the terminal)
 - 2: stderr – where error messages go (usually the terminal)
- `mycmd <infile >outfile 2>&1`
 - Redirect stderr to wherever stdout is currently redirected to
- `grep dyalog /etc/passwd | wc -l`
 - identify all lines in `/etc/passwd` which have `dyalog` in them, and return how many that is
- `echo $?`
 - Print the exit code from the last command run. 0 indicates success, anything else failure of some sort

UNIX cheat sheet – processes

- ps
 - ps -ef list all processes with useful information about each
 - ps -elf list full information about all processes
 - ps -fu andys list all the processes belonging to andys
- top/topas/nmon (AIX)
 - Show most active processes. Ctrl-c stops this
- kill
 - kill -2 pid signal a weak interrupt to process with process id pid
 - kill -3 pid signal a strong interrupt to process with process id pid
 - kill -9 pid unconditionally and immediately kill process with process id pid
 - This should be a last resort, not a first resort !

UNIX cheat sheet – shell/environment variables

ANDY=1	# defines a shell variable, only visible to this shell
export ANDY	# is now available to all commands started from this # shell. An Environment Variable
export ANDY=1	# both the above in one command
echo \$ANDY	# Must put a \$ in front of a variable to reference it

UNIX cheat sheet – watch out for quotes

```
echo "echo $ANDY"      # Substitute value of variable  
echo 1
```

```
echo 'echo $ANDY'      # exactly as was typed  
echo $ANDY
```

```
echo `echo $ANDY`      # treat as though a command  
1
```

UNIX cheat sheet – common environment vars

`$HOME` # my home directory

`$PATH` # list of directories to look for commands in

`$LANG` # defines your language environment

`$TZ` # defines your time zone

`$TERM` # defines what terminal type you are using

- Normally has the value `xterm`
- Get this wrong and you may have an unusable terminal session

□sh and)sh

- Run a command from within APL
 - AIX: !!! time taken to call □SH proportional to size of calling process !!!
 - Avoid using □SH wherever possible

```
        )sh echo $PPID
123456
        □sh 'grep 620 /etc/passwd'
jenkins:!:6203:6116:~/home/jenkins:/usr/bin/bash
nicolas:!:6204:6116:~/home/nicolas:/usr/bin/bash
        +□sh 'grep 620 /etc/passwd'
jenkins:!:6203:6116:~/home/jenkins:/usr/bin/bash nicolas:!:6204:6116:~/home/nicolas:/usr/bin/bash
```

- Quite intentional !

When it goes horribly wrong

- Dyalog APL can "crash"
 - Be careful when using the word crash
 - Do you mean the APL code has errored, but you're still in APL ?
 - Do you mean the APL interpreter has crashed ?
- Types of abnormal termination
 - Syserror - generates an aplcore
 - Core dump - generates a core file if O/S is configured to do so
 - Process dump - need to configure Windows to do so
 - Silently disappears
 - Are you sure the APL code didn't call `⌵OFF` ?

When it goes horribly wrong

- Your tty session may be unusable
 - Close it down and open a new one
 - Ctrl-j stty sane ctrl-j may get control back
- With a syserror you may be able to see where the problem lies with
sed -n '/===== Interesting Information/, \$p' aplcore
- You may be able to) copy or □c y from an aplcore
 -) copy copies global value, □CY the most local

IBM 3151 ASCII terminal








Terminal Emulator

- Windows
 - PuTTY
 - Needs IME and APL-385 font
 - KEA Attachmate
 - Customer configured keyboard
- Linux
 - Xterm and other terminal windows
 - Generally Unicode

Classic issues

- Can't use \subseteq    \ddot{o} $\ddot{ö}$ ı
 - No slots free in □av (nor some fonts)
 - Write cover functions/operators and use them
- Many classic workspaces still use the underscored alphabet
 - Problems rendering them in terminal emulators
 - Lack of keystrokes available
 - May not paste well between APL and other applications

Default keyboard layouts

- Classic / Windows
 - Various different layouts over the years, last updated with 12.1
- Unicode / Windows
 - Uses IME
- Unicode/Linux terminal emulators
 - Supplied with recent Linux distributions
- Classic / PuTTY
 - Uses the Unicode layout & IME
- Classic/KEA
 - Ctrl-o Ctrl-n, underscored alphabet

Keyboard commands

- Listed in UNIX User Guide, Appendix A
 - http://help.dyalog.com/17.1/#UNIX_IUG/Appendix%20A%20Table%20of%20keycodes-te.htm
- Classic: use `␣KL` to see what the keystroke is
 - Need to know the 2 character command mnemonic
 - Can have multiple keystrokes for one command
 - Will be in 18.2 for Unicode

Configuration parameters (Dyalog env vars)

- Are environment variables or command line parameters
- On UNIX MUST be in uppercase
- From 17.1 can appear in config files in \$HOME/.dyalog

MAXWS=1G mapl	# Start APL with a 1G workspace
APLK=kea_apl mapl	# Set my input translate table to kea_apl

APLK/APLKEYS, APLT/APLTRANS

- Classic only
- APLK/APLKEYS
 - Defines what keystroke is used to enter each character
 - Defines what keystroke is used to enter each keyboard command
- APLT/APLTRANS
 - Defines what the interpreter must send to generate each character
 - Defines how to generate colours
- APLKEYS/APLTRANS are similar to PATH: a list of directories to look in
- APLK/APLK0, APLT/APLT1/APLT2 is the name of the file to use
 - Generally is the same as the value of \$TERM

Useful URLs

- help.dyalog.com/17.1 and docs.dyalog.com/17.1
 - UNIX-specific documentation
- www.dyalog.com/user-meetings/index.htm
 - Uncle Andy's Fireside chats (and The Doctor is in)
- www.ibm.com/docs/en/aix/7.2?topic=commands
 - On most UNIX boxes,
 - `command -?`
 - `command -h`
 - `man command`